

Introduction to EDA with PySpark

This presentation will guide you through an Exploratory Data Analysis (EDA) on the NYC Taxi Trip dataset using the powerful PySpark library. We'll dive into the data, uncover insights, and learn how to leverage PySpark's capabilities for effective data analysis.



Problem Statement

The problem statement of the NYC Taxi Data Analysis and Optimization project is to analyze the vast amount of taxi trip data from New York City (NYC) and identify opportunities to optimize various aspects of the taxi system. This includes improving operational efficiency, enhancing service quality, and maximizing revenue for taxi drivers and companies.

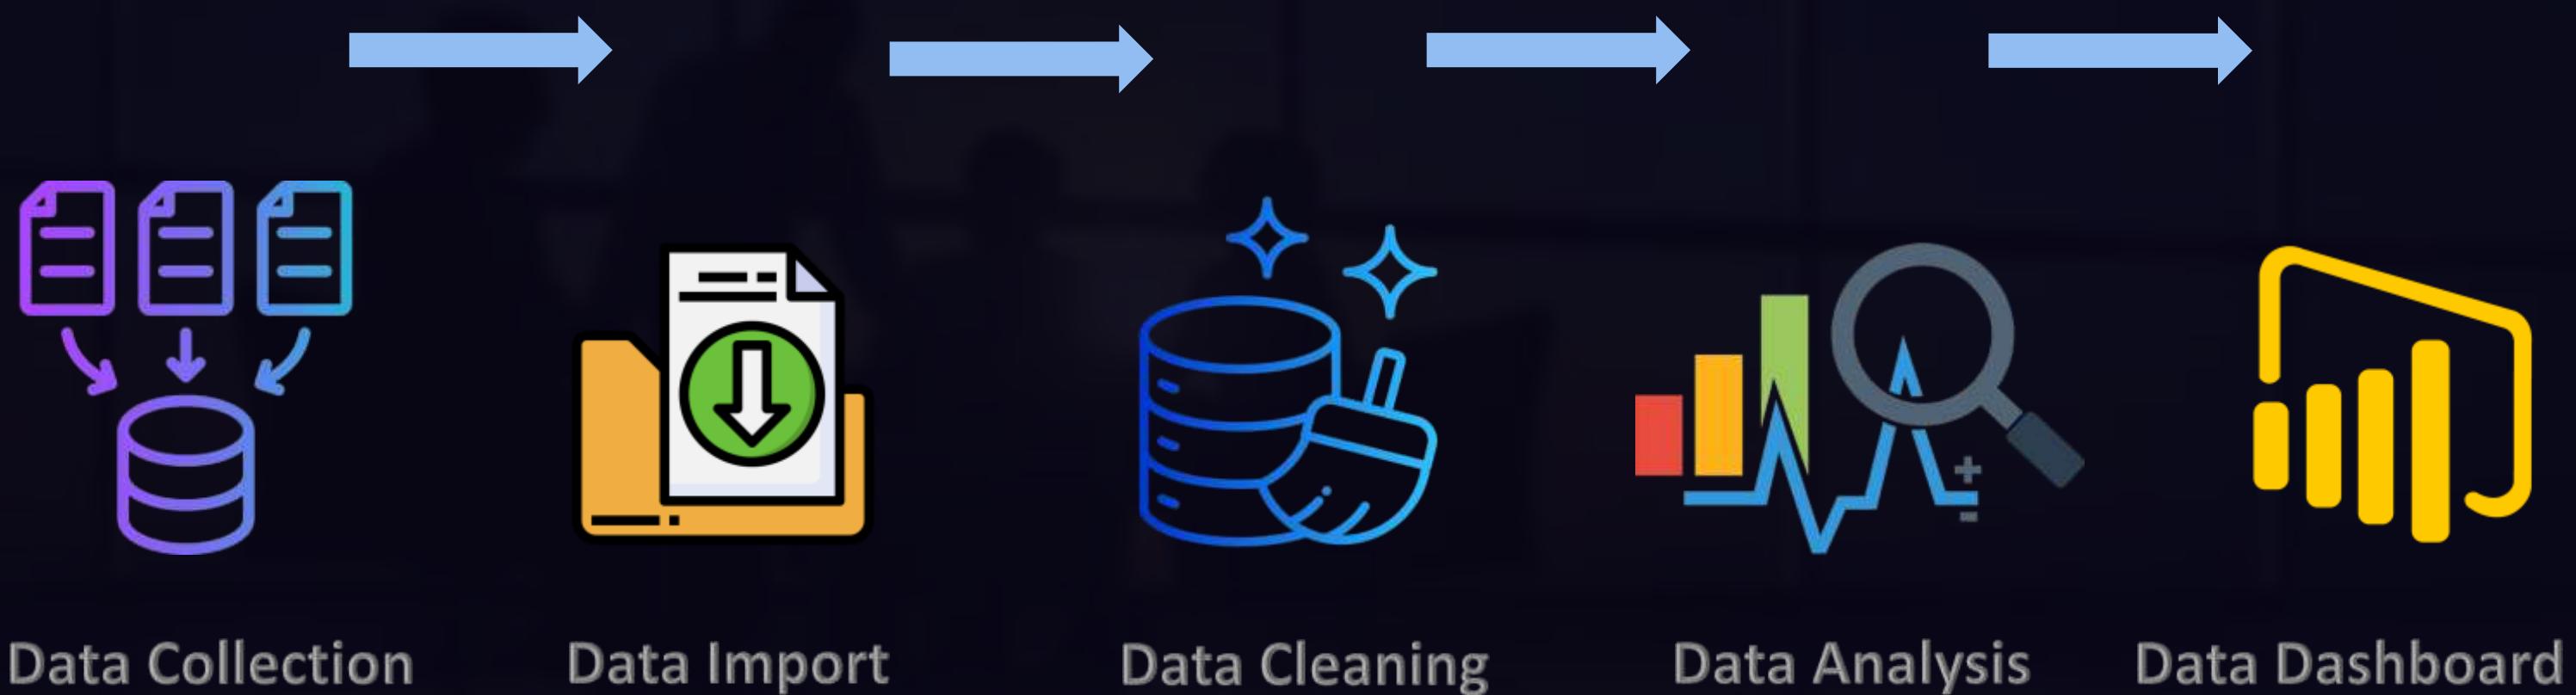
- What are the pickup and drop hotspots?
- Peak Hours to book a trip
- Average Passengers per trip
- Most Trip type within or intercity
- Toll Price
- Most busy days of a week
- Average Tip amount
- How distance change the price of the ride?

These are some the questions we need answer in order to increase revenue, manage the taxi availability and overall estimate the profit from each trip.

So let's get started.



Workflow of the Project



Tools for the Project



Overview of NYC Taxi Trip Data

1

Vendor ID

A code indicating the TPEP provider that provided the record.

1= Creative Mobile Technologies, LLC 2= VeriFone Inc.

2

RateCodeID

The final rate code in effect at the end of the trip.

1= Standard rate 2=JFK 3=Newark 4=Nassau or Westchester 5=Negotiated fare 6=Group ride

3

Payment_type

A numeric code signifying how the passenger paid for the trip.

1= Credit card 2= Cash 3= No charge 4= Dispute 5= Unknown 6= Voided trip



Overview of NYC Taxi Trip Data

Columns: ['VendorID', 'lpep_pickup_datetime', 'lpep_dropoff_datetime', 'store_and_fwd_flag', 'RatecodeID', 'PULocationID', 'DOLocationID', 'passenger_count', 'trip_distance', 'fare_amount', 'extra', 'mta_tax', 'tip_amount', 'tolls_amount', 'ehail_fee', 'improvement_surcharge', 'total_amount', 'payment_type', 'trip_type', 'congestion_surcharge']

1 Vendor ID

A code indicating the TPEP provider that provided the record.

1= Creative Mobile Technologies, LLC 2= VeriFone Inc.

2 RateCodeID

The final rate code in effect at the end of the trip.

1= Standard rate 2=JFK 3=Newark 4=Nassau or Westchester 5=Negotiated fare 6=Group ride

3 Payment_type

A numeric code signifying how the passenger paid for the trip.

1= Credit card 2= Cash 3= No charge 4= Dispute 5= Unknown 6= Voided trip

4 Trip_type

A numeric code signifies type of trip.

1= City Trip 2= InterCity Trip



Importing the Dataset

1

Collect The Data

We collected the data from Kaggle and store it in our local disk as a csv file and description of the data as a txt file

2

Import the data to Databricks' database

We open out Databricks profile. Go to catalog choose a database and create table. Then upload the data from the local storage and read the first row as Header.

3

Read the data in `spark.read.csv`

Create a spark data frame as spark.read.csv and give the file path to the table we just created

	VendorID	lpep_pickup_datetime	lpep_dropoff_datetime	store_and_fwd_flag	RatecodeID	PULocationID	DOLocationID	passenger_count	trip_distance	fare_amount	extra	mta_tax	tip_amount	tolls_amount	improvement_surcharge	Total_amount
0	1	2021-07-01 00:30:52	2021-07-01 00:35:36	N	1	74	168	1	1.20	6	0.5	0.5				
1	2	2021-07-01 00:25:38	2021-07-01 01:01:31	N	1	116	265	2	13.89	42	0.5	0.5				
2	2	2021-07-01 00:05:58	2021-07-01 00:12:00	N	1	97	33	1	.95	6.5	0.5	0.5				
3	2	2021-07-01 00:41:40	2021-07-01 00:47:23	N	1	74	42	1	1.24	6.5	0.5	0.5				
4	2	2021-07-01 00:51:32	2021-07-01 00:58:46	N	1	42	244	1	1.10	7	0.5	0.5				
—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
83686	None	2021-07-02 07:59:00	2021-07-02 08:33:00	None	None	218	169	None	18.04	50.24	2.75	0				
83687	None	2021-07-02 07:02:00	2021-07-02 07:18:00	None	None	74	137	None	5.56	19.16	0	0				
83688	None	2021-07-02 07:53:00	2021-07-02 08:15:00	None	None	69	75	None	5.13	22.45	0	0				
83689	None	2021-07-02 07:56:00	2021-07-02 08:30:00	None	None	117	82	None	12.58	48.62	2.75	0				
83690	None	2021-07-02 07:00:00	2021-07-02 07:26:00	None	None	218	196	None	11.32	45.84	2.75	0				

Handling Missing Values and Outliers

1

Identifying the Columns with Null Values

We'll thoroughly examine the dataset for missing values and understand their potential impact on our analysis.

2

Dropping the columns with 90% Null Values

Depending on the percentage of missing values in a column, we are filtering the dataset.

3

Replacing the None Values with the most common or mean value of the column

We are filling the missing values with the mean or mode of the column based on categorical or numerical value.



dataricks

csv->Parquet Python

```
File Edit View Run Help Last edit was 19 minutes ago New cell UI: OFF
```

Run all nm6 Publish

```
1 # PREPROCESSING THE DATA
2 from pyspark.sql.functions import col
3 for i in df.columns:
4     print(f"{i}: {df.filter(col(i).isNull()).count()}")
5 print(df.count())

▶ (4) Spark Jobs
VendorID: 32518
lpep_pickup_datetime: 0
lpep_dropoff_datetime: 0
store_and_fwd_flag: 32518
RatecodeID: 32518
PUlocationID: 0
DOLocationID: 0
passenger_count: 32518
trip_distance: 0
fare_amount: 0
extra: 0
mta_tax: 0
tip_amount: 0
tolls_amount: 0
improvement_surcharge: 0
total_amount: 0
payment_type: 32518
trip_type: 32518
congestion_surcharge: 32518
83691
Command took 0.62 seconds -- by rohitgo853@gmail.com at 11/04/2024, 10:28:40 on nmE
```

Cmd 5

dataricks

csv->Parquet Python

```
File Edit View Run Help Last edit was 40 minutes ago New cell UI: OFF
```

Run all nm6 Share Publish

```
83691
Command took 0.62 seconds -- by rohitgo853@gmail.com at 11/04/2024, 14:28:46 on nmE
```

Cmd 5

```
1 # Drop columns with 100% null values
2 n=df.count()
3 for i in df.columns:
4     x=df.filter(col(i).isNull()).count()
5     if (x/n)*100 == 100:
6         df = df.drop(i)

▶ (33) Spark Jobs
83691
Command took 6.79 seconds -- by rohitgo853@gmail.com at 11/04/2024, 18:16:16 on nmE
```

Cmd 6

```
1 for i in df.columns:
2     x=df.select(mean(i).alias("nn")).head()[0]
3     y=df.select(mode(i).alias("nn")).head()[0]
4     if df.filter(col(i).isNull()).count()>0:
5         df= df.fillna({i:y})
```

dataricks

csv->Parquet Python

```
File Edit View Run Help Last edit was 40 minutes ago New cell UI: OFF
```

Run all nm6 Share Publish

```
83691
Command took 9.62 seconds -- by rohitgo853@gmail.com at 11/04/2024, 14:28:46 on nmE
```

Cmd 5

```
1 # Drop columns with 100% null values
2 n=df.count()
3 for i in df.columns:
4     x=df.filter(col(i).isNull()).count()
5     if (x/n)*100 == 100:
6         df = df.drop(i)

▶ (33) Spark Jobs
83691
Command took 6.79 seconds -- by rohitgo853@gmail.com at 11/04/2024, 18:16:16 on nmE
```

Cmd 6

```
1 for i in df.columns:
2     x=df.select(mean(i).alias("nn")).head()[0]
3     y=df.select(mode(i).alias("nn")).head()[0]
4     if df.filter(col(i).isNull()).count()>0:
5         df= df.fillna({i:y})
```

dataricks

csv->Parquet Python

```
File Edit View Run Help Last edit was 41 minutes ago New cell UI: OFF
```

Run all nm6 Share Publish

```
1 for i in df.columns:
2     print(f"{i}: {df.filter(col(i).isNull()).count()}")

▶ (31) Spark Jobs
VendorID: 0
lpep_pickup_datetime: 0
lpep_dropoff_datetime: 0
store_and_fwd_flag: 0
RatecodeID: 0
PUlocationID: 0
DOLocationID: 0
passenger_count: 0
trip_distance: 0
fare_amount: 0
extra: 0
mta_tax: 0
tip_amount: 0
tolls_amount: 0
improvement_surcharge: 0
total_amount: 0
payment_type: 0
trip_type: 0
congestion_surcharge: 0
```

How many trips Recorded?

83691

```
Python ▶▶ ▾ - ×
```

```
1 # How many trips Recorded?
2 df.count()
```

▶ (2) Spark Jobs

Out[187]: 83691

Command took 0.49 seconds -- by rohitgo853@gmail.com at 11/04/2024, 14:28:46 on nm6

#What is the average trip distance?

194.3546993105587
7

```
1 #What is the average trip distance?
2 df.select(mean("trip_distance").alias("Average_trip_distance")).show()
3
```

▶ (2) Spark Jobs

Average_trip_distance
194.35469931055877

What are the top 5 PULocationIDs where trips start from?

PULocationID	count
74	8770
75	7713
41	4761
42	3229
95	2486

```
Cmd 13

1 # What are the top 5 PULocationIDs where trips start from?
2 df.groupBy("PULocationID").count().orderBy(col("count").desc()).limit(5).show()

▶ (2) Spark Jobs

+-----+-----+
|PULocationID|count|
+-----+-----+
|      74| 8770|
|      75| 7713|
|      41| 4761|
|      42| 3229|
|      95| 2486|
+-----+-----+
```

#What are the top 5 DOLocationIDs where trips end?

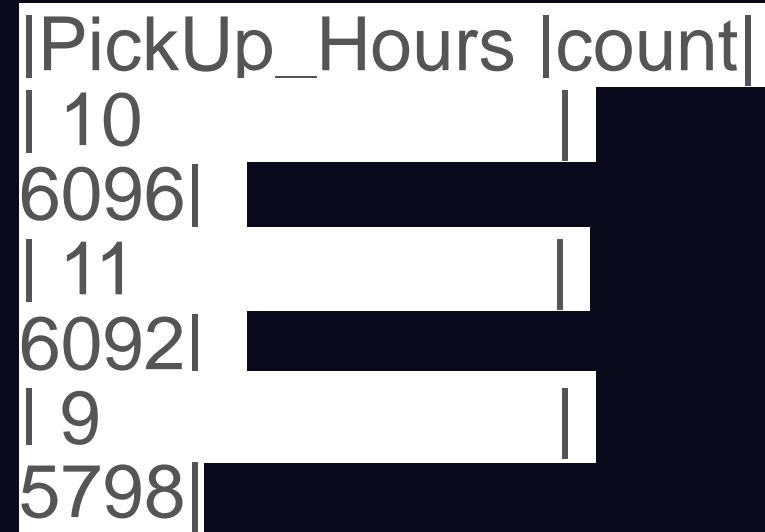
DOLocationID	count
74	3666
75	3122
42	2904
41	2527
236	1700

```
1 #What are the top 5 DOLocationIDs where trips end?
2 df.groupBy("DOLocationID").count().orderBy(col("count").desc()).limit(5).show()
3

▶ (2) Spark Jobs

+-----+-----+
|DOLocationID|count|
+-----+-----+
|      74| 3666|
|      75| 3122|
|      42| 2904|
|      41| 2527|
|     236| 1700|
+-----+-----+
```

Peak Hours for Someone to start a ride

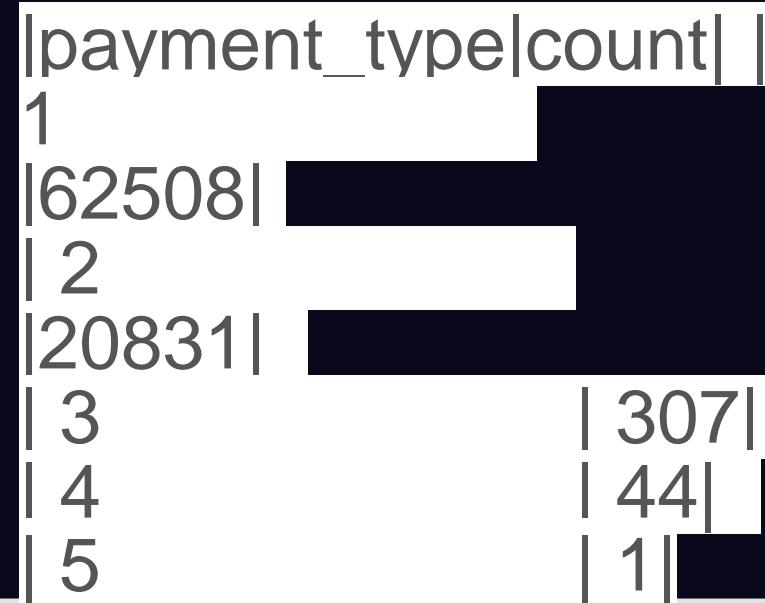


```
1 # Peak Hours for Someone to start a ride
2 df.groupBy(hour("lpep_pickup_datetime").alias("PickUp_Hours")).count().orderBy(col("count").desc()).limit(5).show()
```

▶ (2) Spark Jobs

```
+-----+-----+
|PickUp_Hours|count|
+-----+-----+
| 10| 6096|
| 11| 6092|
| 9| 5798|
| 12| 5766|
| 15| 5744|
+-----+-----+
```

What is the distribution of trips based on the payment type?



```
1 # What is the distribution of trips based on the payment type?
2 df.groupBy("payment_type").count().orderBy("payment_type").show()
3
```

▶ (2) Spark Jobs

```
+-----+-----+
|payment_type|count|
+-----+-----+
| 1|62508|
| 2|20831|
| 3| 307|
| 4| 44|
| 5| 1|
+-----+-----+
```

What is the most common RatecodeID in the dataset?

RatecodeID count
1 1
81512 [REDACTED]
5 1954 [REDACTED]
2 158 [REDACTED]
4 41 [REDACTED]
3 26 [REDACTED]

```
Cmd 17

1 # What is the most common RatecodeID in the dataset?
2 df.groupby("RatecodeID").count().orderBy(col("count").desc()).limit(10).show()

▶ (2) Spark Jobs
+-----+-----+
|RatecodeID|count|
+-----+-----+
| 1| 1 |
| 81512| [REDACTED] |
| 5| 1954 |
| 2| 158 |
| 4| 41 |
| 3| 26 |
+-----+-----+
```

Average fare amount for different vendors

1 17.013627772674567
[REDACTED]
2 24.925072694291224

```
1 # Average fare amount for different vendors
2 df = df.withColumn("total_amount", col("total_amount").cast(DoubleType()))
3 df.groupby("VendorID").mean("total_amount").show()

▶ (2) Spark Jobs
▶ df: pyspark.sql.dataframe.DataFrame = [VendorID: string, lpep_pickup_datetime: string ... 17 more fields]
+-----+-----+
|VendorID| avg(total_amount)|
+-----+-----+
| 1|17.013627772674567 |
| 2|24.925072694291224 |
+-----+-----+
```

#What is the distribution of passenger counts?

```
| 1|76645|  
| 2| 3922|  
| 5| 1240|  
| 6| 1018|  
| 3| 626|  
| 4| 181|
```

```
1 #What is the distribution of passenger counts?  
2 df.groupBy(col("passenger_count").alias("No. of passengers")).count().orderBy(col("count").desc()).show()  
  
▶ (2) Spark Jobs  
  
+-----+-----+  
|No. of passengers|count|  
+-----+-----+  
| 1|76645|  
| 2| 3922|  
| 5| 1240|  
| 6| 1018|  
| 3| 626|  
| 4| 181|  
| 0| 56|  
| 7| 2|  
| 32| 1|  
+-----+-----+
```

#What is the average fare amount for trips?

20.38830459667026

```
1 #What is the average fare amount for trips?  
2 df.select(mean("fare_amount")).show()  
  
▶ (2) Spark Jobs  
  
+-----+  
| avg(fare_amount)|  
+-----+  
|20.38830459667026|  
+-----+
```

How many trips have a tolls amount greater than zero?

Average Tolls
6.413921953613792

Tolls_amount greater than zero
8149

```
1 # How many trips have a tolls amount greater than zero?
2 df.filter(col("tolls_amount")>0).count()
3 n=df.filter(col("tolls_amount")>0)
4 n.select(mean("tolls_amount").alias("Average Tolls_amount"),count("tolls_amount").alias("Tolls_amount greater than zero")).show()
```

▶ (4) Spark Jobs

▶ █ n: pyspark.sql.dataframe.DataFrame = [VendorID: string, lpep_pickup_datetime: string ... 17 more fields]

```
+-----+-----+
|Average Tolls_amount|Tolls_amount greater than zero|
+-----+-----+
| 6.413921953613792|                 8149|
+-----+-----+
```

What is the distribution of trip types?

| City Trip|81931|
| InterCity Trip| 1760|

```
1 # What is the distribution of trip types?
2 df= df.withColumn("TRIPS", when(df["trip_type"]==1,"City Trip")
3           .when(df["trip_type"]==2,"InterCity Trip"))
4
5 df.groupBy(col("TRIPS").alias("No. of passengers")).count().orderBy(col("count").desc()).show()
```

▶ (2) Spark Jobs

▶ █ df: pyspark.sql.dataframe.DataFrame = [VendorID: string, lpep_pickup_datetime: string ... 18 more fields]

```
+-----+-----+
|No. of passengers|count|
+-----+-----+
|      City Trip|81931|
| InterCity Trip| 1760|
+-----+-----+
```

```
# What is the most common day of the week for trips
```

```
| Sat|14964|  
| Fri|14929|  
| Sun|12862|  
| Thu|11752|  
| Wed|10875|  
| Tue|10403|  
| Mon| 7906|
```

```
1 # What is the most common day of the week for trips?  
2 df=df.withColumn("DATE",col("lpep_pickup_datetime").cast("date")).withColumn("DayW",dayofweek(col("lpep_pickup_datetime")))  
3  
4 df=df.withColumn("DW",when(df["DayW"]==1,"Mon")  
5 .when(df["DayW"]==2,"Tue")  
6 .when(df["DayW"]==3,"Wed")  
7 .when(df["DayW"]==4,"Thu")  
8 .when(df["DayW"]==5,"Fri")  
9 .when(df["DayW"]==6,"Sat")  
10 .when(df["DayW"]==7,"Sun"))  
11  
12 df.groupby(col("DW").alias("day of week")).count().orderBy(col("count").desc()).show()  
13 df=df.drop(col("DATE"),col("DW"),col("DayW"))  
  
▶ (2) Spark Jobs  
▶ ⏷ df: pyspark.sql.DataFrame = [VendorID: string, lpep_pickup_datetime: string ... 10 more fields]  
  
Day of Week|count|  
---|---|  
Sat|14964|  
Fri|14929|  
Sun|12862|  
Thu|11752|  
Wed|10875|  
Tue|10403|  
Mon| 7906|
```

```
# Is there any correlation between trip distance and total amount paid?
```

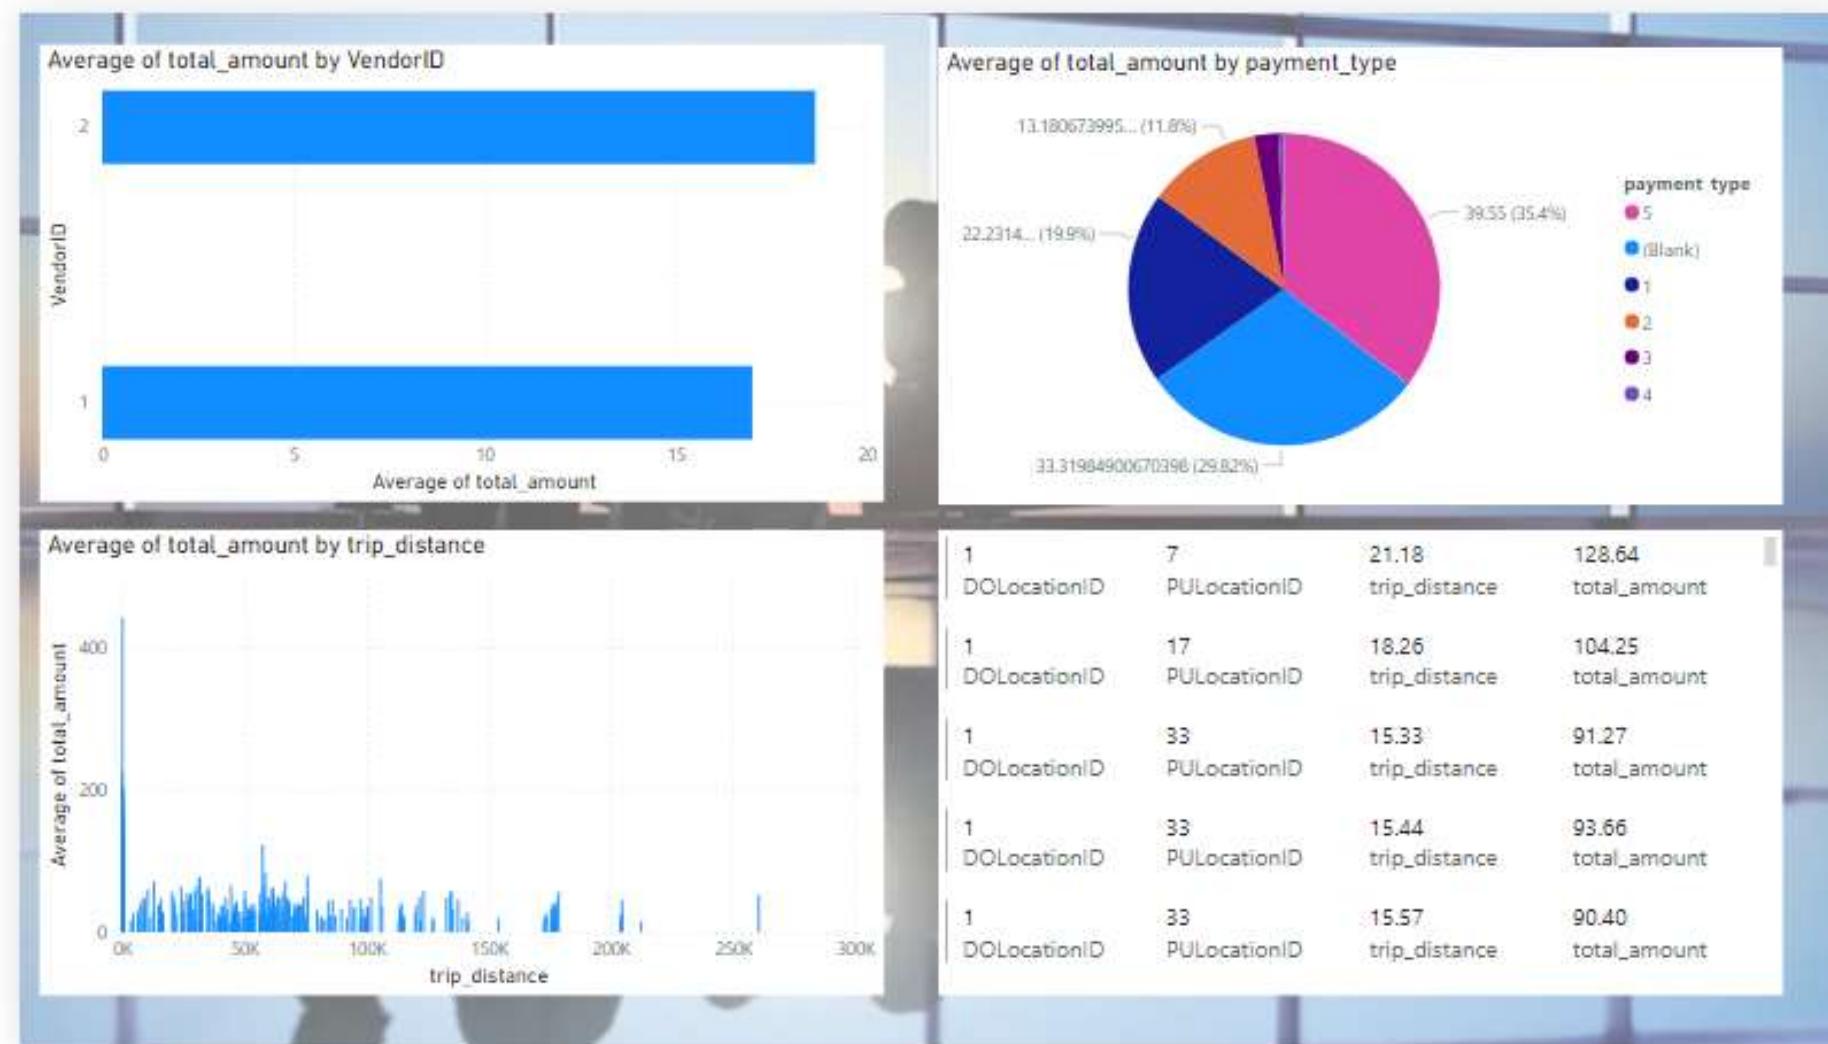
```
0.02507182116693  
05
```

```
1 # Is there any correlation between trip distance and total amount paid?  
2 df.select(corr("trip_distance","total_amount")).first()[0]
```

```
▶ (2) Spark Jobs
```

```
Out[251]: 0.0250718211669305
```

Data Dashboard in PowerBI



Conclusion

- Pickup Location ID hotspots are 74, 75, 41, 42, 95
- Drop Location ID hotspots are 74, 75, 42, 236
- Peak Hours for a ride is 9-12 am and 3pm
- Trips are mostly pay in credit cards and cash
- Average fare amount of Vendor 1 is 17.01 and Vendor 2 is 24.92
- Mostly number of passengers for a ride is 1, 2 or 5
- Passengers mostly use within City Trip_type
- Average Toll amount for 8149 rides is 6.41
- Saturday and Friday are the most busy days of the week
- Average tip amount is 1.05
- The relation between trip distance and total amount is weak positive 0.025.

